

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [2]: quality = pd.read_csv("Wine_Quality.csv")
```

```
In [3]: quality.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

So in here we have 11 independent features and 1 dependent feature(quality).

```
In [6]: quality.shape
```

Out[6]: (1599, 12)


```
In [7]: quality.dropna(inplace = True)
```

```
In [8]: quality.describe()
```

Out[8]:


	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1596.000000	1596.000000	1596.000000	1596.000000	1596.000000	1596.000000	1596.000000
mean	8.321366	0.527666	0.271128	2.536936	0.087487	15.882206	46.43107
std	1.742121	0.179154	0.194847	1.408341	0.047107	10.467380	32.89307
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

```
In [9]: plt.figure(figsize = (12,6))  
sns.countplot(x=quality['quality'])  
plt.show()
```




Overall, this code creates a countplot to visualize the distribution of values in the 'quality' column of the DataFrame. It helps in understanding the frequency or occurrence of different quality levels and provides insights into the data distribution.

```
In [10]: plt.figure(figsize = (12,6))  
sns.barplot(x='quality', y = 'alcohol', data = quality, palette = 'inferno')  
plt.show()
```

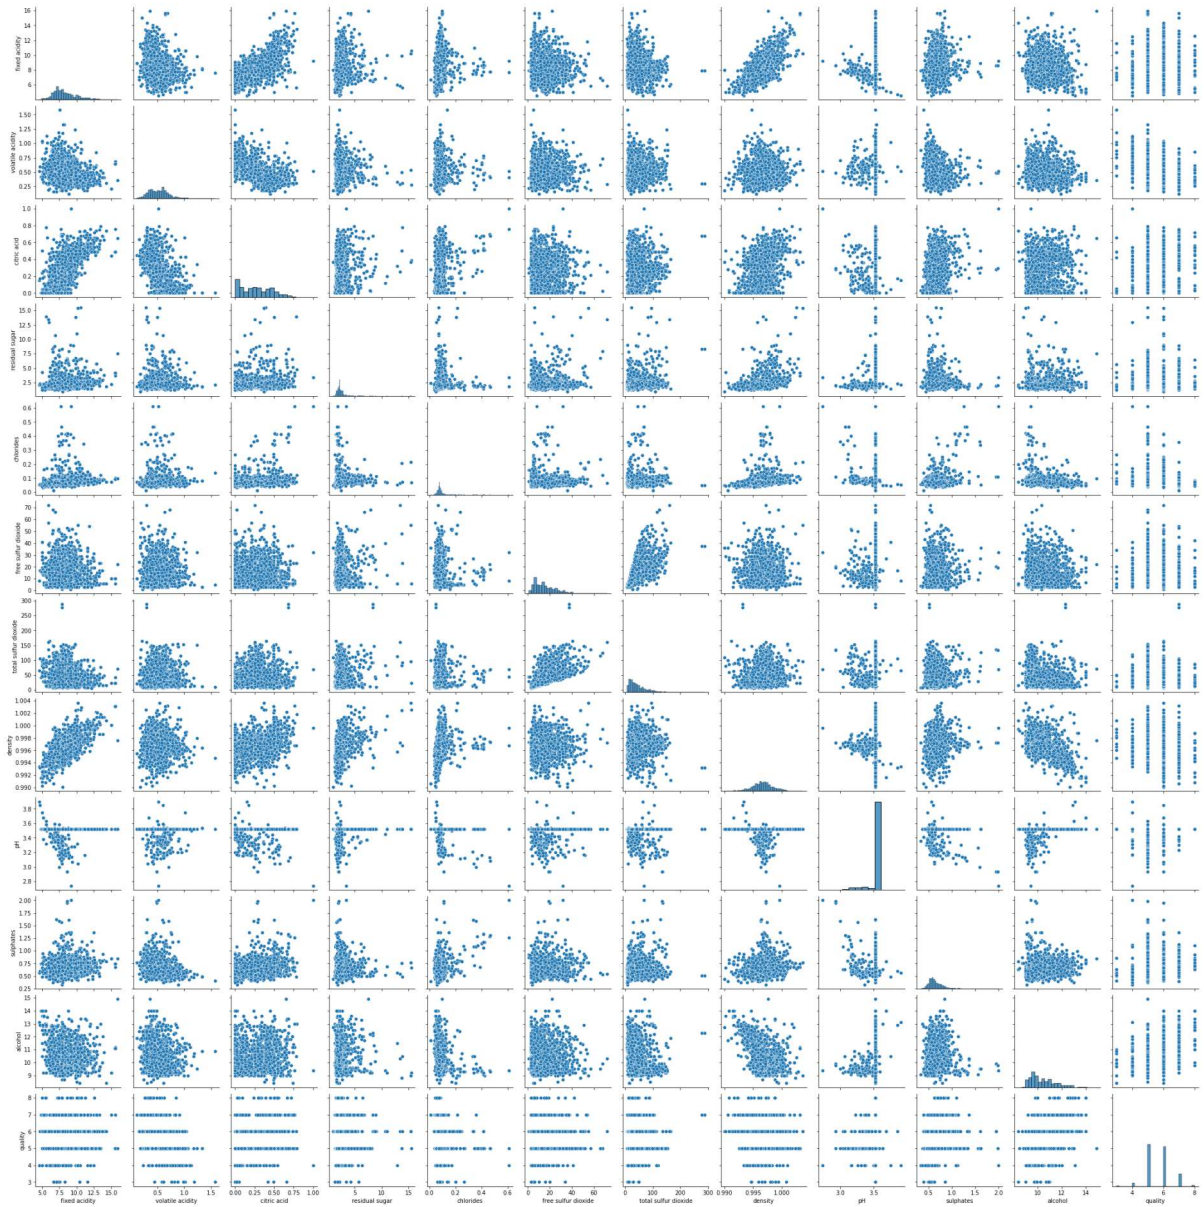


```
In [11]: plt.figure(figsize = (12,6))  
sns.scatterplot(x='citric acid', y = 'pH', data = quality)  
plt.show()
```




```
In [12]: plt.figure(figsize = (12,6))  
sns.pairplot(quality)  
plt.show()
```

<Figure size 864x432 with 0 Axes>



```
In [13]: plt.figure(figsize = (12,6))
sns.heatmap(quality.corr())
plt.show()
```



```
In [14]: x=quality.drop(['quality'], axis=1)
y=quality['quality']
```

```
In [15]: from sklearn.ensemble import IsolationForest
```

```
In [16]: # Outlier Detection using Isolation Forest
outlier_detector = IsolationForest(contamination=0.05)
outlier_labels = outlier_detector.fit_predict(x)
```

C:\Users\user\anaconda3\lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but IsolationForest was fitted with feature names
warnings.warn(

```
In [17]: # Remove outliers from the dataset
x = x[outlier_labels == 1]
y = y[outlier_labels == 1]
```

```
In [ ]:
```

In summary, the code separates the features (independent variables) into the DataFrame x by excluding the 'quality' column, and assigns the target variable (dependent variable) to the Series y by extracting only the 'quality' column.

This kind of separation is commonly done when preparing data for machine learning models, where the features and the target variable need to be handled separately.

Splitting the data into dependent and independent variables before preprocessing ensures that the preprocessing steps are applied consistently to both sets.

If you perform preprocessing for entire dataset before splitting, there is a risk of data leakage, where information from the test set or target variable leaks into the training set, leading to biased results.

Preprocessing steps such as scaling or normalization are typically applied to the features (independent variables) rather than the target variable (dependent variable).

Data Preprocessing

In [18]: `pip install imbalanced-learn`

Requirement already satisfied: imbalanced-learn in c:\users\user\anaconda3\lib\site-packages (0.11.0)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: numpy>=1.17.3 in c:\users\user\anaconda3\lib\site-packages (from imbalanced-learn) (1.21.5)
Requirement already satisfied: scipy>=1.5.0 in c:\users\user\anaconda3\lib\site-packages (from imbalanced-learn) (1.7.3)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\user\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
Requirement already satisfied: joblib>=1.1.1 in c:\users\user\anaconda3\lib\site-packages (from imbalanced-learn) (1.2.0)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\user\anaconda3\lib\site-packages (from imbalanced-learn) (1.2.2)

In [19]: `## oversampling
from imblearn.over_sampling import SMOTE

os = SMOTE()
x_res, y_res = os.fit_resample(x, y)`

The `fit_sample()` method performs oversampling by generating synthetic samples from the minority class (es) in the dataset, in order to balance the class distribution. It creates new samples by interpolating between existing samples of the minority class. The resulting `x_res` and `y_res` are the oversampled feature matrix and target variable, respectively.

Oversampling is typically used when dealing with imbalanced datasets, where one or more classes

are underrepresented compared to other classes. In such cases, oversampling helps to address the class imbalance by creating synthetic samples of the minority class(es). This can improve the performance of machine learning models by providing a more balanced representation of the classes during training.

```
In [20]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_res,y_res,test_size=0.2,
```

```
In [21]: from sklearn.preprocessing import StandardScaler

stdscale = StandardScaler().fit(x_train)
x_train_std = stdscale.transform(x_train)
x_test_std = stdscale.transform(x_test)
```

```
In [22]: from sklearn.metrics import accuracy_score
```

Logistic Regression

```
In [23]: from sklearn.linear_model import LogisticRegression

# Train a Logistic regression model
lr = LogisticRegression()
lr.fit(x_train_std, y_train)

# Make predictions on the test set
predictions = lr.predict(x_test_std)

accuracy_score(y_test, predictions)
```

```
Out[23]: 0.5613577023498695
```

```
In [28]: from sklearn.metrics import classification_report
```

```
In [29]: # Evaluate model performance
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
3.0	0.76	0.97	0.85	117
4.0	0.57	0.53	0.55	124
5.0	0.47	0.44	0.46	123
6.0	0.43	0.28	0.34	143
7.0	0.49	0.50	0.50	137
8.0	0.58	0.72	0.64	122
accuracy			0.56	766
macro avg	0.55	0.57	0.56	766
weighted avg	0.54	0.56	0.55	766

Decision Tree Classifier

```
In [30]: from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()
dt.fit(x_train_std, y_train)

accuracy_score(y_test, dt.predict(x_test_std))
```

Out[30]: 0.7924281984334204

Random Forest Classifier

```
In [31]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state = 42)
rf.fit(x_train_std, y_train)

accuracy_score(y_test, rf.predict(x_test_std))
```

Out[31]: 0.835509138381201

```
# Hence Random Forest Classifier is the best classification model to predict
the wine test quality as
it has the highest accuracy score.
```